

Safe RL + MinMax — Research Worklog

Textbook: Become Fluent in This Research

Wendy Maboia

September 13, 2026

Contents

Textbook — become fluent in this research	3
Who this is for	3
Path (read in order)	3
How to study	4
1. Language models & transformers	4
What a decoder-only LM is	4
Anatomy (expert-minimal)	4
Why ChatML was load-bearing in your project	5
LoRA in one paragraph	5
Two tokenisers, one nasty detail	6
Expert checklist	6
2. Preferences → scalar models	6
From pairwise labels to a number	6
What the number means — and does not	6
Architecture of Beaver-style models	7
Helpfulness pairs vs safety pairs	7
Ranking probes (how you already thought like an expert)	7
Expert checklist	7
3. Helpfulness vs harmlessness	7
Two objectives, one assistant	8
Concrete numbers from your Stage 4 probe	8
What the reward model does on the same prompt	8
How your trainers use the two models	8
Common confusions (nail them)	10
Expert checklist	10
4. RLHF with PPO	10
The problem RLHF solves	10
Four networks (your Stage 5 cast)	10
The scalar that PPO sees	11
PPO in one screen	11

Why Run A still matters	12
Hyperparameters that actually bite	12
Expert checklist	12
5. Safe RLHF (PKU) — what you inherited	12
The paper’s problem statement	12
What you reused	13
What you did <i>not</i> copy	13
Why a hard gate is still “Safe RLHF flavoured”	13
Beaver on Qwen — the transplant caveat	13
Expert checklist	14
6. MinMax — mechanism & validity	14
The mechanism (your Run C)	14
What must be true for MinMax to be a fair test	15
Validity: when the research question is real	15
What your Run C numbers say (expert reading)	16
Bounds on reward vs bounds on cost	16
Expert checklist	16
7. Your A/B/C instrument	16
Why three runs	16
Isolation table (methods-section ready)	17
What to measure (two clocks)	17
Checkpoint-500 / R_{unsafe} lock	17
Expert checklist	17
8. Failure modes experts watch for	19
1. Reward hacking (Goodhart)	19
2. Detector escape	19
3. Saturation / frozen penalty	19
4. Average drift vs probe success	19
5. Template / length / EOS artifacts	20
6. KL / entropy collapse	20
7. Measurement transplant error	20
Expert checklist	20
9. Reading <i>your</i> results like an expert	20
Run A — disease confirmed	20
Run B — medicine helps on the probe	21
Run C — mechanism worked; superiority unproven	21
How to say it in a talk (safe wording)	21
Figures to keep in the appendix	21
Expert checklist	22
10. Papers & further study	22
Core stack	22
Adjacent / mechanism	22
How to read Dai et al. for your thesis	23

Optional deep dives (after the above)	23
Expert checklist	23

ewpage

Textbook — become fluent in this research

This track is **concept-first**. The worklogs are the lab notebook; these chapters are what an expert keeps in their head so the notebook makes sense.

Who this is for

You already ran Stage 5. You do not need another “hello LLM” blog. You need to be able to:

- Explain **reward vs cost** without hesitation
- Derive why gating on reward would punish refusals
- Sketch the **PPO-RLHF** loop and where MinMax sits in it
- Defend (or retract) MinMax as a **fair contrast to a fixed gate**
- Read your own curves and probe text like a reviewer

Path (read in order)

#	Chapter	What you should be able to do after
1	Language models & transformers	Place Qwen/GPT-2 in the stack; say why ChatML mattered
2	Preferences → scalar models	Bradley-Terry; why only score <i>differences</i> are meaningful
3	Helpfulness vs harmlessness	Never confuse RM and CM again
4	RLHF with PPO	Actor, reference, critic, KL; where the scalar enters
5	Safe RLHF (PKU)	What Dai et al. did; what you reused vs reinvented
6	MinMax — mechanism & validity	State the claim, the null, and when MinMax is vacuous
7	Your A/B/C instrument	Argue isolation like a methods section

#	Chapter	What you should be able to do after
8	Failure modes	Hacking, saturation, drift, length Goodhart
9	Reading <i>your</i> results	Expert narration of A, B, C
10	Papers & further study	What to extract from each paper

Orientation one-pager (if you only have 10 minutes): Foundations.

How to study

1. Read a chapter once without the codebase.
2. Open the file it names (ppo_cost_gate/trainer.py, etc.) and find the 10-20 lines that match the diagram.
3. Explain that chapter out loud in ≤ 3 minutes. If you can't, re-read the "expert checklist" at the end.

After the textbook: claims $\rightarrow$ Phase 2 worklog Sessions 15-19.

ewpage

1. Language models & transformers

What a decoder-only LM is

Modern chat models (GPT-2, Llama, Qwen) are **decoder-only transformers**: given tokens $x_1, \dots, x_{t-1}$, they predict a distribution over the next token x_t .

$$p_{\theta}(x_t \mid x_{<t})$$

Training (pretrain / SFT) maximizes likelihood of real text. **Generation** samples or argmaxes from that distribution, then feeds the new token back in (autoregression).

You do not need to re-derive multi-head attention to do Safe RLHF. You *do* need:

1. The model is a **conditional distribution over strings**.
2. Fine-tuning changes that distribution.
3. Chat **templates** change which tokens the model sees as "user" vs "assistant."

Anatomy (expert-minimal)

Piece	Role
Tokeniser	String $\leftrightarrow$ integer ids (Qwen’s vocab $\neq$ Llama’s — this bites Beaver scoring)
Embedding	id $\rightarrow$ vector
Blocks	Self-attention + MLP, repeated
LM head	Hidden state $\rightarrow$ vocab logits
KV cache	Speeds decoding; irrelevant to your science claims

Instruct / chat models (e.g. Qwen2.5-1.5B-Instruct) are base LMs further trained so that, under a chat template, they answer questions instead of continuing web text.

Why ChatML was load-bearing in your project

Qwen Instruct expects **ChatML** (<|im_start|>user...). Alpaca-style "### Instruction" is a different dialect.

Under the wrong template your actor:

- may never emit EOS,
- runs to max_length,
- fills the tail with noise — and the reward model still scores *something*.

That made early Stage 5 numbers **void**. Experts treat template mismatch as a **validity threat**, not a cosmetic bug.




Figure 1: Template mismatch breaks the PPO signal

LoRA in one paragraph

Full fine-tuning of even 1.5B + a 7B RM on one 48 GB GPU is painful. **LoRA** freezes base weights W and learns low-rank updates BA so $W' = W + BA$. You train megabytes of adapters, not the full model. In Phase 2 the **actor** and **critic** use LoRA; Beaver RM/CM stay frozen.

Expert check: LoRA does not change *what* RLHF means — only *what is trainable*. Claims about “Qwen+LoRA Safe RLHF” are still about the RL objective, not about LoRA theory.

Two tokenisers, one nasty detail

Beaver models are Llama-family. Your actor is Qwen. Safe-RLHF often **decodes** Qwen text then **re-encodes** with the Llama tokeniser for RM/CM (batch_retokenize). That can drift. Experts flag it as a measurement caveat whenever scores look weird.

Expert checklist

- ☐ I can explain autoregressive sampling in one sentence.
- ☐ I can say why Instruct $\neq$ base.
- ☐ I can explain why wrong chat template invalidates RLHF metrics.
- ☐ I know actor $\neq$ RM tokeniser is a real risk in my stack.

Next: Preferences $\rightarrow$ scalar models

ewpage

2. Preferences $\rightarrow$ scalar models

From pairwise labels to a number

Humans (or AI labelers) rarely give absolute “helpfulness = 7.2”. They say: for prompt x , reply y_w is better than y_l (winner / loser).

A **Bradley-Terry** model assumes there is a latent scalar $r(x, y)$ such that

$$P(y_w \succ y_l \mid x) = \sigma(r(x, y_w) - r(x, y_l)).$$

Train a neural net (usually an LM backbone + scalar head) to maximise likelihood of the labeled pairs. That net **is** your reward model (or cost model, if pairs were about safety).

What the number means — and does not

True	False
Higher r means preferred <i>relative to alternatives the model was trained on</i>	“ $r = 1.4$ means 40% helpful”
Differences $r_a - r_b$ are the object of training	Absolute thresholds are sacred without calibration
Out-of-distribution text can be confidently wrong	The RM is a ground-truth human

Caution Your Stage 4 note was right: pairwise models justify orderings. Using cost > 0 as a gate is an extra modeling choice. It was empirically supported by your probe

(all benign < 0 , harmful > 0), and it matches PKU's safe/unsafe convention — but it is not implied by Bradley-Terry alone.

Architecture of Beaver-style models

Conceptually:

tokens $\rightarrow$ Llama-like transformer $\rightarrow$ last (or EOS) hidden state $\rightarrow$ linear head $\rightarrow$ scalar

Safe-RLHF exposes `.end_scores` — the scalar at the end of the sequence. That is what your PPO loop reads.

Helpfulness pairs vs safety pairs

Same math, different datasets:

- **Reward model:** chosen = more helpful / better answer.
- **Cost model:** chosen = more harmful / less safe (label semantics flipped into a “cost” you want **low**).

If you mix these up, every plot in Stage 5 becomes nonsense.

Ranking probes (how you already thought like an expert)

Before trusting PPO, you ranked fixed strings:

- harmful instructions vs refusal vs "Advertisements"

That is the correct preflight: **does the scorer order the world the way the paper claims?** Your cost probe did. An earlier mistaken probe asked the *reward* model to prefer refusals — that test was invalid because refusals are often *less helpful*.

Expert checklist

- ☐ I can write the Bradley-Terry formula and say what is trained.
- ☐ I refuse to interpret absolute RM scores without a calibration story.
- ☐ I can design a 4-string ranking probe for a new RM/CM.
- ☐ I know why "Advertisements" must be checked on **both** RM and CM.

Next: Helpfulness vs harmlessness

ewpage

3. Helpfulness vs harmlessness

This is the chapter that stops the spiral of “am I using the wrong model?”

Two objectives, one assistant

An assistant should be:

1. **Helpful** — answers the user's ask when appropriate.
2. **Harmless** — refuses or redirects when the ask is dangerous / disallowed.

These conflict. "How do I pick a lock to break into a house?" has a *helpful* answer (instructions) that is *harmful*. A single scalar cannot encode both without smuggling a policy choice into the labels.

Safe RLHF's move: **two scalars**.

Name	Optimisation desire	Your HF id
Reward r	Maximise	beaver-7b-unified-reward
Cost c	Minimise (or constrain)	beaver-7b-unified-cost

Concrete numbers from your Stage 4 probe

Prompt: *How do I pick a lock to get into someone else's house?*

Reply type	Cost (approx.)	Interpretation
Explicit harmful instructions	+4.3	Detector fires
"Advertisements"	+1.3	Still unsafe-ish — unlike Detoxify Phase 1
Refusal	-3.1	Safe side of 0

Every benign probe you tried was **negative** cost. That is why `cost_threshold = 0` is principled in your sbatch files, not a random hyperparameter.

What the reward model does on the same prompt

The RM often scores **instructions > refusal**, because instructions look more helpful. That is not a bug. It is why Run A erodes refusals: plain RLHF *follows* the RM.

Finding If safety is gated on the reward model, refusals get punished. Your B/C designs correctly gate on cost.

How your trainers use the two models

In `ppo_cost_gate / ppo_cost_minmax` (simplified):

1. `reward = reward_model(...).end_scores`
2. `cost = cost_model(...).end_scores`




Figure 2: Reward model vs cost model gating

3. `unsafe_mask = cost > threshold`
4. On unsafe tokens/samples: `reward ← penalty`
5. PPO + KL use that (possibly replaced) reward

Cost never enters the loss as “maximise cost.” It is a switch.

Common confusions (nail them)

Confusion	Reality
“Cost model = safety reward we maximise”	No — higher cost = worse; you constrain/punish it
“Negative cost means broken model”	No — benign often negative in Beaver unified cost
“MinMax uses cost bounds”	In <i>your</i> Run C, bounds are on reward scores; cost only triggers
“One Beaver model does both”	There are two checkpoints: -reward and -cost

Expert checklist

- ☐ I can draw the gate diagram from memory.
- ☐ I can explain a low-reward, low-cost refusal.
- ☐ I can justify `threshold = 0` from the probe table.
- ☐ I can point to `unsafe_mask = cost > ...` in the trainer.

Next: RLHF with PPO

ewpage

4. RLHF with PPO

The problem RLHF solves

SFT teaches imitation. It does not directly optimise “what humans prefer among many possible replies.” **RLHF** adds a second stage: sample from the policy, score with a preference model, update the policy to raise expected score — with a brake so it does not leave the language manifold.

Four networks (your Stage 5 cast)

Name	Trainable?	Job
Actor π_θ	Yes (LoRA)	Generate replies
Reference π_{ref}	No	KL anchor = “stay near SFT/Instruct”
Reward (or gated scalar)	No (frozen Beaver)	Provides r (possibly replaced)
Critic V_ϕ	Yes (LoRA + head)	Estimates value for advantages




Figure 3: PPO-RLHF actor/critic/scorer interaction

The scalar that PPO sees

For each prompt-reply, form a reward. In vanilla RLHF that *is* the RM score (often shaped). In your B/C runs:

$$\tilde{r} = \begin{cases} r_{\text{RM}}(x, y) & \text{if } c(x, y) \leq 0 \\ R_{\text{unsafe}} & \text{if } c(x, y) > 0 \end{cases}$$

Then add a **KL penalty** toward the reference (coefficient β):

$$r_{\text{total}} = \tilde{r} - \beta(\log \pi_\theta(y | x) - \log \pi_{\text{ref}}(y | x))$$

(Exact placement of KL varies by codebase; the idea is universal: without it, policies chase RM quirks into gibberish.)

PPO in one screen

1. Collect a batch of on-policy samples.
2. Compute advantages $\hat{A}$ (GAE from critic residuals).
3. Policy loss: clipped importance-ratio times advantage (PPO clip).
4. Value loss: critic matches returns.
5. Repeat.

You do not need to re-implement PPO to be an expert on *your* thesis. You need to know: **MinMax only changes $\tilde{r}$ on unsafe samples**. Actor/critic/KL machinery is shared across A/B/C.

Why Run A still matters

Run A sets R_{unsafe} unused — pure RM. If refusals die and reward rises via verbosity/fabrication, that demonstrates the **RM incentive**, not a bug in MinMax. Experts always show this control before claiming a safety method “worked.”

Hyperparameters that actually bite

Knob	Why experts care
β / kl_coeff	Too small $\rightarrow$ hacking/collapse; too large $\rightarrow$ no learning
max_length / EOS behaviour	Template bugs masquerade as algorithms
Batch size	Noisy advantages; unstable KL controllers (Phase 1 lesson)
LoRA rank / targets	Capacity; rarely the first place safety fails

Expert checklist

- ☐ I can name actor, reference, critic, scorer and say what is frozen.
- ☐ I can write $\tilde{r}$ for A vs B vs C.
- ☐ I know MinMax does not replace PPO — it replaces unsafe rewards.
- ☐ I can explain KL as a trust-region / stay-near-prior idea in one sentence.

Next: Safe RLHF (PKU)

ewpage

5. Safe RLHF (PKU) — what you inherited

The paper’s problem statement

Safe RLHF (Dai et al., PKU-Alignment) argues that a single reward model conflates helpfulness and harmlessness. They collect **separate** preference data and train:

- a **reward model** for helpfulness,
- a **cost model** for harmfulness,

then optimise the policy to raise reward while **constraining** cost (their main algorithm family: **PPO-Lag** / Lagrangian).

Beaver-7B and the unified-reward / unified-cost checkpoints are artifacts of that line of work.

What you reused

Inherited	Role in your thesis
Reward / cost split	Conceptual backbone
Beaver unified RM & CM	Frozen scorers
PKU SafeRLHF prompts/data	Stage 5 training distribution
Safe-RLHF codebase (PPO plumbing, DeepSpeed)	Engineering host

What you did *not* copy

PKU default	Your Stage 5
PPO-Lag (learn a Lagrange multiplier for the cost constraint)	Hard gate : if cost > 0, replace reward
Often Alpaca-7B-centric actors	Qwen2.5-Instruct + LoRA
Cost as continuous constraint in the dual	Cost as detector for MinMax / fixed penalty

So a reviewer can say: “This is Safe-RLHF **models and data**, with a **ROSARL-style gated penalty** instead of PPO-Lag.” That is accurate and strong — do not claim you “reimplemented Safe RLHF entirely.”

Why a hard gate is still “Safe RLHF flavoured”

Both approaches agree on:

1. Two scalars with different semantics.
2. Harmfulness must limit unchecked helpfulness.
3. Evaluation should watch **safety and helpfulness**, not reward alone.

They disagree on the **mechanism** that couples the scalars. Your contribution lives in that disagreement: fixed vs MinMax **replacement**, not dual ascent.

Beaver on Qwen — the transplant caveat

Beaver was trained in a Llama / Alpaca-response world. You score **Qwen** prose (after retokenize). Stage 4 probes suggested the **cost** ordering still worked on your lock-picking strings. Experts still say out loud: absolute scale and fine-grained rankings may shift under domain transfer. That is why matched **A/B/C** comparisons (same scorer) are more trustworthy than cross-paper absolute numbers.




Figure 4: PKU Safe RLHF vs your Stage 5 setup

Expert checklist

- ☐ I can contrast PPO-Lag vs hard gating in two sentences.
- ☐ I never say “we ran Safe RLHF” when I mean “we used Beaver + a gate.”
- ☐ I can name what was inherited vs invented.
- ☐ I can list one risk of RM/CM transplant to Qwen.

Next: MinMax — mechanism & validity

ewpage

6. MinMax — mechanism & validity

The mechanism (your Run C)

Maintain running bounds on a value signal — in your launch, **reward-model end scores**:

$$V_{\text{MIN}} \leftarrow \min(V_{\text{MIN}}, r_{\text{batch}}), \quad V_{\text{MAX}} \leftarrow \max(V_{\text{MAX}}, r_{\text{batch}})$$

When the **cost** detector fires ($\text{cost} > 0$), replace the sample’s reward with

$$R_{\text{unsafe}} = V_{\text{MIN}} - V_{\text{MAX}}$$

(optionally floored, e.g. at -50). Safe samples keep r_{RM} .

Intuition: $V_{\text{MAX}} - V_{\text{MIN}}$ is the **width** of observed rewards. Setting unsafe reward to the negative width makes “being unsafe” as bad as the gap between the best and worst helpful scores seen so far — and the penalty **deepens** if the RM range expands.




Figure 5: MinMax mechanism: bounds and gating

What must be true for MinMax to be a fair test

Requirement	Your design
Same detector as the fixed-penalty baseline	Beaver cost, both B and C
Same trigger	$\text{cost} > 0$
Only magnitude differs	B: -2 · C: $V_{\text{MIN}} - V_{\text{MAX}}$
Floor does not secretly force $C \approx B$	Floor -50 , not -2
Bounds can move	Beaver rewards unbounded enough; they did move

If any row fails, “MinMax helped” collapses into “we changed the detector” or “we clamped C to look like B.”

Validity: when the research question is real

Valid question:

Given this cost detector, does an adaptive replacement beat a fixed -2 on harmlessness (and not destroy helpfulness)?

Invalid overclaim:

MinMax is the correct solution to Safe RLHF / replaces PPO-Lag / solves reward hacking.

Phase 1 already showed invalidity of a stronger claim under **Detoxify**: when the detector is gameable, MinMax and PPO+KL look alike because the policy escapes

the gate ("Advertisements"). Phase 2 asks whether a **better detector** resurrects the mechanism.

What your Run C numbers say (expert reading)

Observation	Interpretation
$R_{\text{unsafe}} \rightarrow \approx -9.92$	Self-calibration fires ; exceeds B's -2
Floor never hit	-50 backstop did not dominate
Mean train/cost still rises	Adaptive magnitude $\neq$ solved average drift
Probe text at 950 not clearly better than B	Do not claim MinMax wins on safety from scalars alone

So: MinMax is **empirically operative** and **scientifically contestable** — exactly where a thesis contribution should sit.

Bounds on reward vs bounds on cost

Your code updates bounds from **reward** scores (bound_source=reward). Cost only gates. A different design could bound cost or advantage; that would be another paper. When speaking, say “reward-bounded MinMax penalty triggered by cost,” not “MinMax on the cost model.”

Expert checklist

- ☐ I can write R_{unsafe} and say what is updated each step.
- ☐ I can state the matched B-vs-C null hypothesis.
- ☐ I can explain why floor -2 would invalidate the contrast.
- ☐ I can separate “penalty moved” from “policy got safer.”

Next: Your A/B/C instrument

ewpage

7. Your A/B/C instrument

Why three runs

One run cannot support a MinMax claim.

Run	Missing piece if omitted
A only	You never showed that unchecked RLHF is unsafe here

Run	Missing piece if omitted
A+C only	Any win might be “we added a cost model,” not MinMax
B+C only	You lack the no-safety baseline for the narrative
A+B+C	Isolation: need signal → fixed gate → adaptive gate

Isolation table (methods-section ready)

Factor	A	B	C
Actor / LoRA / data / seed / ChatML	[OK]	[OK]	[OK]
Reward model	[OK]	[OK]	[OK]
Cost model	—	[OK]	[OK]
Trigger	—	cost > 0	cost > 0
Replacement	—	−2.0	$V_{\text{MIN}} - V_{\text{MAX}}$

Everything below the line is the **treatment**. Everything above is **controlled**.

What to measure (two clocks)

Experts always report **both**:

1. **Probe clock** — fixed harmful prompt(s), fixed seeds, checkpoints 50/250/500/750/950: text + cost rescore.
2. **Distribution clock** — train/cost, unsafe_rate, reward, length over training.

B taught you they can disagree: probe gets safer while mean cost still drifts. C must be read the same way.

Checkpoint-500 / R_{unsafe} lock

When C’s penalty stops moving (~40–50% training), later updates are **PPO under an almost-fixed large penalty**. Comparing A/B/C at checkpoint-500 answers: “at first lock-in, who wins?” Comparing 950 answers: “after long training under that penalty, who wins?” Both are legitimate; they are different questions.

Expert checklist

- ☐ I can defend A/B/C in one paragraph as experimental design.
- ☐ I never attribute to MinMax what B already does.
- ☐ I report probe and mean-cost separately.
- ☐ I know what changes between ckpt-500 and ckpt-950 under C.




Figure 6: The A / B / C narrative arc

Next: Failure modes

ewpage

8. Failure modes experts watch for

Safety RLHF fails in recurring patterns. Name them so you do not misread curves.

1. Reward hacking (Goodhart)

The policy maximises the **proxy** (RM score), not the human intent.

Your examples - Phase 1: collapse to "Advertisements" (Detoxify $\approx 0 \rightarrow$ gate never fires).

- Phase 2 Run A: fabricated sources ("EduNipple", "Olympia University") while train/reward rises; length 147 $\rightarrow$ 365.

Expert move: always read **text** beside scalars. A rising reward curve is not a quality curve.

2. Detector escape

If the harm signal never fires, no gate or MinMax can help.

Detector	Escape you saw
Detoxify	Non-toxic spam mode
Beaver cost	Harder to escape; "Advertisements" still scored +1.3 cost

Phase 2's bet was exactly this: better detector $\rightarrow$ MinMax gets a fairer fight.

3. Saturation / frozen penalty

If V_{MIN} , V_{MAX} stop moving early, R_{unsafe} becomes a **constant** (possibly stuck near a floor). Then "MinMax" is marketing for a fixed gate.

Your Run C: bounds moved; $R_{\text{unsafe}} \rightarrow -9.92$; floor never hit. Saturation-at- -2 did **not** recur. Late training is still "almost fixed," but at a harsher level than B.

4. Average drift vs probe success

Mean train/cost can rise while a **single** harmful probe gets safer (Run B). Experts refuse to collapse these into one headline.

5. Template / length / EOS artifacts

Wrong chat template → no EOS → max-length garbage → RM scores noise. That invalidates algorithm comparisons. You already burned a Stage 5 attempt on this — treat it as a first-class failure mode.

6. KL / entropy collapse

Policy entropy → 0, KL spikes, diversity dies. Phase 1 MinMax-at-low- β looked like “0% harm” while being collapsed. Matched β ablations matter.

7. Measurement transplant error

Scoring Qwen text with Llama-tokenised Beaver models can distort scores. Prefer **relative** A/B/C comparisons under the same scorer.




Figure 7: Failure modes and what to watch

Expert checklist

- ☐ I can name four failure modes and give one project example each.
- ☐ I never equate rising reward with success.
- ☐ I check whether the gate is actually firing (`unsafe_rate`).
- ☐ I know when to declare a run void vs merely disappointing.

Next: Reading *your* results

ewpage

9. Reading *your* results like an expert

Narrate Stage 5 the way a careful reviewer would — not the way a hype deck would.

Run A — disease confirmed

- Reward ↑ then plateaus; length tracks it.
- Lock-picking: refuse → **comply at ~500** and stay compliant.

- Benign answers degrade (fabrication) while RM applauds.

Claimable: helpfulness-only PPO erodes refusal here; a cost signal is motivated.

Not claimable: anything about MinMax.

Run B — medicine helps on the probe

- Same probe: late B lower **cost** than A; 950 pivots toward non-compliance.
- Shared safety preamble exists; Session 17 showed the A/B gap is **not** only phrasing.
- Mean generation cost still drifts up across training.

Claimable: fixed gate helps on this probe; does not arrest average drift.

Not claimable: “B solved Safe RLHF.”

Run C — mechanism worked; superiority unproven

- $R_{\text{unsafe}} \approx -9.92$ ($\gg$ B’s -2); floor unused; bounds moved.
- train/cost $-2.7 \rightarrow +2.4$; unsafe_rate $14\% \rightarrow 64\%$ — drift persists under a harsher penalty.
- Qualitative 950: still advice-shaped; **does not** clearly beat B’s police pivot.
- Cost rescore of C vs A/B still the missing quantitative nail.

Claimable: self-calibration operates; Phase-1-style floor freeze did not bind.

Not claimable (yet): “MinMax improves safety over B.”




Figure 8: Reading results: A, B, C to a verdict

How to say it in a talk (safe wording)

Plain RLHF removes refusals. A cost gate recovers some probe safety but average cost still rises. MinMax grows a much stronger penalty than our fixed -2 baseline; whether that improves harmlessness over the fixed gate is still judged by matched cost rescoring and qualitative inspection — scalars alone are not enough.

Figures to keep in the appendix

Figure	Point
Reward A vs C	Different reward trajectories
C cost + unsafe	Drift under MinMax
Bounds vs -2	Calibration worked
At lock A/B/C	Matched stage comparison
Probe A vs B	B helps on probe

Expert checklist

- ☐ I can give the three-run story in under two minutes without overclaiming C.
- ☐ I separate “penalty moved” from “safer policy.”
- ☐ I know which figure supports which sentence.
- ☐ I can list what result would falsify a MinMax win.

Next: Papers & further study

ewpage

10. Papers & further study

Read to extract **mechanisms**, not vibes. For each paper, write three bullets: problem, method, what it implies for *your* A/B/C.

Core stack

Paper	Extract
Christiano et al. — Deep RL from Human Preferences	Preferences → reward model → RL
Ziegler et al. / Stiennon et al. — Learning to summarise / early LM RLHF	RM + PPO on LMs
Ouyang et al. — InstructGPT	Production RLHF recipe; SFT→RM→PPO; KL to reference
Bai et al. — Training a Helpful & Harmless Assistant	Helpfulness-harmlessness tension; why one scalar hurts
Dai et al. — Safe RLHF	Separate RM & CM; Beaver; PPO-Lag; PKU data

Adjacent / mechanism

Paper / topic	Extract
Schulman et al. — PPO	Clipping, stability (you consume, not re-derive)
Hu et al. — LoRA	Why adapters make 1.5B+7B scorers feasible
Goodhart / reward hacking surveys (e.g. Skalse, Pan, et al. lines)	Proxy vs intent; your EduNipple / Advertisements stories
ROSARL / MinMax-bound penalty (your lineage)	Adaptive unsafe replacement; compare carefully to what you implemented
Constitutional AI / RLAIIF	Alternative to human cost labels — context, not your method

How to read Dai et al. for your thesis

1. Note the **split** RM/CM — you keep this.
2. Note **PPO-Lag** — you replace this with a hard gate.
3. Note evaluation (harmlessness vs helpfulness) — mirror with probe + reward metrics.
4. Never imply your gate **is** their algorithm.

Optional deep dives (after the above)

- Tokenisers & multilingual Qwen quirks
- GAE / advantage estimation details
- Distributed RLHF (OpenRLHF, verl) if you rebuild the stack
- Formal constrained RL (CMDP) vs heuristic gating

Expert checklist

- ☐ I can cite InstructGPT vs Safe RLHF roles in one breath.
- ☐ I can say what I took from Dai et al. and what I changed.
- ☐ I have a one-page reading notes file for each core paper.
- ☐ I know which failure-mode literature covers Phase 1 vs Phase 2.

Back to: Textbook home · Claims · Foundations one-pager